

The background of the entire image is a dark blue field filled with a pattern of red dots. These dots are arranged in a way that they form a large, faint circular shape in the center, with the density of the dots being higher in the center and tapering off towards the edges. The dots are of varying sizes, creating a textured, digital effect.

HUST

ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

ONE LOVE. ONE FUTURE.



ĐẠI HỌC
BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY
OF SCIENCE AND TECHNOLOGY

REAL-TIME PAYMENT FRAUD DETECTION FOR AN E-COMMERCE PLATFORM

An End-to-End Business Intelligence Pipeline on PaySim

Authors:

Lương Minh Dương (20251038M) · Trần Lê Phương Thảo (20251186M)

Hoàng Minh Hoàng (20252561M) · Ngô Việt Anh (20252570M)

Phạm Tuấn Kiệt (20252751M) · Nguyễn Như Thái (20252270M)

Group 1 · Business Intelligence (IT5315E) · HUST

ONE LOVE. ONE FUTURE.

1 Business Problem & Cost Model

2 Data & Exploratory Insights

3 Features, Models & Evaluation

4 Deployment, Monitoring & Policy

5 Conclusion



Business Problem & Cost Model

The Problem: Fraud is a Needle in a Haystack

Setting

A payment platform must decide, **at authorization time** (before money moves), whether each transaction is fraudulent - in milliseconds, with only what is known *before* the transaction settles.

Why it is hard

- Fraud is **0.129%** of traffic - a **774:1** imbalance (8,213 of 6.36M).
- A model that flags *nothing* is already **99.87%** accurate - and useless.

So accuracy is the wrong target - the real question is how much money we save, at what operational cost.

Why it matters

- Total fraud exposure: **12.06 billion** currency units.
- A missed fraud costs the *full* amount; a false alarm costs only friction and a little review labour.



Cost-Based Framing, not Accuracy

We minimise an expected cost that turns a three-sided trade-off into one number

$$\text{Cost} = \underbrace{c_{FN} \sum_{i \in FN} \text{amount}_i}_{\text{missed fraud}} + \underbrace{c_{FP} \cdot |FP|}_{\text{friction}} + \underbrace{c_{MR} \cdot |\text{flagged}|}_{\text{review labour}}$$

- $c_{FN} = 1.0 \times \text{amount}$, $c_{FP} = 25$, $c_{MR} = 3$ (kept in `src/config.py`).
- The ratio fixes a **break-even amount of 25**: block whenever the amount at risk exceeds 25 units.

KPIs that pull against each other

- **Loss avoided** (north-star, in money).
- **False-positive rate** (customer friction).
- **Review-queue volume** (a hard capacity limit).





Data & Exploratory Insights


Data: PaySim + a Synthetic Contextual Layer

- **PaySim**: 6,362,620 mobile-money transactions, **8,213 fraud** (0.129%), 743 hourly steps.
- PaySim has *no* device / geo / account context, so we add a **synthetic layer** (Faker): device, IP-billing distance, disposable email, high-risk country, shipping/billing mismatch, night flag.
- Context is attached with `reveal_rate=0.55` and `legit_noise=0.04` - signal is realistic, not perfect.




Class imbalance (log scale): fraud is a thin sliver.

What the Data Reveals (1): Where and How Much



- Fraud appears **only** in **TRANSFER** (0.77%) and **CASH_OUT** (0.18%) - the cash-exit channels.
- Zero fraud in PAYMENT / CASH_IN / DEBIT $\Rightarrow$ we can screen out **~60% of traffic for free**.
- Fraud concentrates in **large amounts**: the top amount decile carries a 0.69% fraud rate.
- **Night** transactions: 1.77% fraud vs 0.10% by day - a behavioural signature of automated attacks.

What the Data Reveals (2): The Leakage Trap



A signal that looks perfect - but cheats

Balance-reconciliation features reach single-feature **AUC 0.92**, yet they are computed from *post-transaction* balances - unknown at auth time.

- Using them is **target leakage** - the score is a *consequence* of fraud, not a predictor.
- The real test: **is this feature available when we must decide?**
- This splits features into **leaky** vs **deployable** groups - our central design choice.

Features, Models & Evaluation

Leakage-aware feature groups


- **Leaky** (*base* 18, *all* 50): post-transaction balances - reference only.
- **Deployable** (*dest* 16, *synth* 21, *realistic* 44): available at auth time.
- Destination history is aggregated **causally** (past transactions only - no future leakage).

Handling 774:1 imbalance

- Class weighting / `scale_pos_weight`, not naive resampling.
- **Time-based split**: train steps 1-323, val 323-378, test 378-743 - train on the past, test on the future.
- Correlations are all $|r| < 0.08$: fraud is a **non-linear** phenomenon.

Model Results & the “Leakage Tax”

Test AUC-PR on deployable features



Model	AUC-PR	Recall	Flagged
LogReg	0.587	0.997	16.8%
Random Forest	0.812	0.953	4.8%
XGBoost	0.907	0.981	2.4%

XGBoost wins: high recall while flagging the *least* traffic. Trees capture the non-linear interactions that a linear model (recall $\approx$ 1 but 16.8% flagged) cannot.

On *leaky* features every model scores AUC-PR $\approx$ 1.0 - the **leakage tax** is the gap down to 0.907.

Cost-selected threshold = 0.241

- Recall **0.981**, precision **0.173**, flagging only **2.38%** of traffic.
- Intercepts **99.91%** of fraud value (loss avoided).

Low precision is **by design**: with a break-even of 25, 5 cheap false alarms are worth 1 six-figure catch. Precision is a *capacity lever*, not a modelling defect.

3-model max-risk ensemble

- Combine LogReg + RF + XGBoost; take the **maximum** risk (most cautious).
- Map to a decision: **block** > **review** > **allow**, with a block floor at 0.9.
- Every decision carries **reason codes** for the review analyst.

Deployment, Monitoring & Policy

Real-time scoring (M6)

- **FastAPI** /score endpoint; single-port in-process option for a fresh clone.
- Returns risk, decision (block/review/allow), and reason codes.
- Review queue prioritised by **expected loss** (amount $\times$ risk), not raw score.

Streamlit operator app

- 5 pages: **Overview, Cost & ROI, Live Feed, Monitoring, API Tester**.
- Turns the pipeline into something an analyst can *operate*, not just a notebook.

Keeping the model honest (M7)

- **PSI drift** per feature; alert when $\text{PSI} > 0.25$.
- One-click / automatic **retraining** on fresh data; version badge + audit log.
- Outbound **webhook alerts** on drift / decay.

The threshold is a business dial, re-tuned as costs and drift change - not a fixed constant.

Risk policy (M8)

- **Auto-block** the highest-risk band (floor 0.9).
- **Human review** the mid band, queued by expected loss.
- **Allow** the low band - keep good customers frictionless.



Conclusion

Conclusion

- We built a **full BI lifecycle** (M1-M8): business framing, EDA, cleaning, leakage-aware features, modelling, deployment, monitoring, and risk policy.
- The winning model - **XGBoost on deployable features** - intercepts **99.91%** of fraud value while flagging only **2.38%** of traffic (recall 0.981).
- The key discipline is **leakage awareness**: an offline AUC of ~ 1.0 is a trap; only features available at authorization time count.

One line

Fraud detection is a **cost-optimisation** problem - its optimum is deliberately high-recall, low-precision, and re-tuned as the business and the data drift.

Team & Role Distribution

Member	Student ID	Role & responsibilities
Lương Minh Dương	20251038M	Data Engineer - data generation, dictionary, cleaning
Trần Lê Phương Thảo	20251186M	Data Analyst - EDA, fraud-pattern visualization, presentation
Hoàng Minh Hoàng	20252561M	ML Engineer - features, classifier, inference pipeline
Ngô Việt Anh	20252570M	ML Engineer - imbalance, cost-based eval, leakage-aware pipeline
Phạm Tuấn Kiệt	20252751M	MLOps Lead - deployment (API + review queue), drift, final report
Nguyễn Như Thái	20252270M	MLOps / Reporting - demo app, monitoring, final report

All code authored by the team; pipeline reproducible from a single fixed seed (SEED=42).

A decorative graphic on the left side of the slide. It features a dark blue background with a large, stylized circular pattern composed of many small red dots. The dots are arranged in concentric, slightly offset rings, creating a sense of depth and movement. The word "HUST" is centered within this pattern.

HUST

THANK YOU !



hust.edu.vn



fb.com/dhbkhn